

# Implementing GMM

March 3, 2022

```
[1]: import numpy as np
from scipy.io import wavfile
import scipy.io
from scipy.stats import multivariate_normal
import matplotlib.pyplot as plt
import os
import math
import random
```

## 0.1 Computing Spectrograms of the audio files

### Observations from analysis of audio files

- All the music files had a sampling frequency of 16K samples per second
- The total number of samples in each audio files was common for all (480000 samples)

Number of samples in 1s = 16K

Number of samples in 25ms =  $16K * 0.025 = 400$  samples

The shift between successive windows =  $10ms$

Number of samples to shift moving from one window to next =  $16K * 0.010 = 160$  samples

**Number of windows** Without adding any amounts of padding at the end so the left end of the window should move only upto  $480000 - 400 = 479600$

We will be shifting by 160 samples each time.

The total number of windows/frames =  $\frac{479600}{160} = 2997.5 = (\text{approx})2998$

```
[2]: class Spectrogram:
    def __init__(self, folder_path=None):
        self.name=folder_path
        self.num_frames = None
        self.data_matrix = None

    def plot_time_series_graph(self, wav_file):
        time = np.linspace(0, length, wav_file.shape[0])
        plt.plot(time, wav_file, label="channel")
        plt.xlabel("time (s)")
```

```

plt.ylabel("Amplitude")
plt.show()

# Computing the number of frames in the spectrogram
def compute_num_frames(self,audiofile):
    starting=0
    increment = 160
    window_size = 400
    self.num_frames = 0
    while(starting+window_size <= len(audiofile)):
        self.num_frames+=1
        starting += increment

# Compute the spectrogram from the data
def compute_spectrogram(self,audiofile):

    starting = 0
    increment = 160
    window_size = 400

    if(self.num_frames is None):
        self.compute_num_frames(audiofile)
    spectrogram = None
    while(starting+window_size < len(audiofile)):
        current_window = np.array(audiofile[starting : starting + ↵
↵window_size])
        #Perform the required transformation
        fft_transform = np.log(np.abs(np.fft.fft(current_window,n=64)[:32] .
↵reshape(32,1))+1e-10)
        if spectrogram is None:
            spectrogram = fft_transform
        else:
            spectrogram = np.append(spectrogram,fft_transform,axis = 0)
        starting += increment
    return spectrogram

# Preparing the Data Matrix for training the size of dataset would be 32 * ↵
↵(Total frames taken from all audio)
def compute_data_matrix(self):

    for file in os.listdir(self.name):
        sample_rate, wav_file = wavfile.read(self.name+"/"+file)
        spectrogram = self.compute_spectrogram(wav_file).reshape((32,-1))
        if( self.data_matrix is None):
            self.data_matrix = spectrogram
        else:
            self.data_matrix = np.append(self.data_matrix,spectrogram,axis=1)

```

```

        #self.plot_spectrogram(spectrogram)
        #print(self.original_data_matrix.shape)

    def plot_spectrogram(self, spectrogram):
        time = np.linspace(0, self.num_frames, num=self.num_frames)
        freq = np.linspace(0, 32, 32)
        plt.pcolormesh(time, freq, spectrogram)
        plt.xlabel("Time")
        plt.ylabel("Frequency")
        plt.show()

```

## 0.2 KMeans Clustering for finding Initial Data centers

```

[3]: # Helper class for KMeans clustering Algorithm
# Reference : https://pythonprogramming.net/k-means-from-scratch-machine-learning-tutorial/

class KMeans:

    def __init__(self, k, data_matrix, max_iter=10, remove_non_diagonal=False):
        self.k = k
        self.n = data_matrix.shape[0]
        self.data_matrix = data_matrix
        self.max_iter = max_iter
        self.remove_non_diagonal = remove_non_diagonal
        self.centers = {}
        self.covs = {}
        self.labels = {}

    # Compute the cluster centers according to the KMeans Algorithm
    def compute_means(self):
        for i in range(self.k):
            self.centers[i] = self.data_matrix[i]

        for i in range(self.max_iter):
            for j in range(self.k):
                self.labels[j] = []
            for p in range(len(self.data_matrix)):
                current_point = self.data_matrix[p]

                # Get the label of the cluster to which the point belongs
                label = np.argmin(np.array([np.linalg.norm(current_point - self.
↪centers[c]) for c in range(self.k)]))
                self.labels[label].append(self.data_matrix[p])

```

```

        # Compute the new centers by taking the mean out of all points
        ↪ assigned to a cluster
        new_centers = {}
        for j in range(self.k):
            mean_vector = np.mean(np.array(self.labels[j]),axis=0)
            new_centers[j] = mean_vector
        self.centers = new_centers

    # Compute the covariance of data points
    # present in each cluster
    # after finding the cluster centers

    def compute_covariance(self):
        for i in range(self.k):
            new_covs = np.cov(np.array(self.labels[i]).transpose())

            # Remove Non Diagonal entries if needed
            if(self.remove_non_diagonal):
                diagonal_elements = new_covs.diagonal()
                new_covariance_matrix = np.eye(new_covs[i].shape[0])
                for j in range(self.k):
                    new_covariance_matrix[j][j] = diagonal_elements[j]
                new_covs = new_covariance_matrix
            self.covs[i] = new_covs

```

## 1 Class for training the GMM Model

In a GMM Model we are interested in the probability

$$P(X_i|\Theta) = \sum_{k=1}^K \alpha_k \mathbb{N}(X_i|\Theta_k)$$

The parameters  $\alpha$  and  $\mu_k, \Sigma_k$  is determined by using the EM Algorithm where in the E step we determine

$$p(Z_i = l|X_i, \Theta_n) = \frac{\alpha_l \mathbb{N}(X_i|Z_i=l, \Theta_n)}{\sum_{k=1}^K \alpha_k \mathbb{N}(X_i|Z_i=l, \Theta_n)}$$

where we assign soft values of membership of data point to a cluster

The M step the parameters are updated such that

$$\alpha_{\{new\}}^{\{l\}} = \frac{1}{N \sum_{n=1}^N P(Z_i=l|X_i, \Theta_n)}$$

$$\mu_l^{new} = \frac{\sum_{n=1}^N X_i P(Z_i=l|X_i, \Theta_n)}{\sum_{n=1}^N P(Z_i=l|X_i, \Theta_n)}$$

$$\mu_l^{new} = \frac{\sum_{n=1}^N P(Z_i=l|X_i, \Theta_n) (X_i - \mu_l^{new})(X_i - \mu_l^{new})^T}{\sum_{n=1}^N P(Z_i=l|X_i, \Theta_n)}$$

Here the initial values of the mean and covariances are determined using KMeans Algorithm  
alpha is taken to be  $1/K$  where  $K$  is the number of mixture components

```
[4]: class GMM:

    def __init__(self, K, data, skip_non_diagonal_entries=False, iterations=10):
        self.K = K
        self.means = None
        self.covs = None
        self.alpha = {}
        for i in range(self.K):
            self.alpha[i] = 1.0/K

        self.clip_nondiag_cov = skip_non_diagonal_entries
        self.data = data
        self.num_iterations = iterations
        self.likelihoods = None
        self.plot_points = []

        # Initialize the mean covariance matrices of the clusters
        def initialize(self, centers, cov_matrix):
            self.means = centers
            self.covs = cov_matrix

        # Computing the posterior probability of belonging to a cluster given a
        ↪ datapoint
        # and model parameters
        def compute_posterior(self, X):

            probability = np.zeros((X.shape[0], self.K))

            # Denominator term
            for i in range(self.K):
                normal_fn = multivariate_normal(mean=self.means[i], cov=self.covs[i])
                for j in range(X.shape[0]):
                    pdf_value = normal_fn.pdf(X[j])
                    probability[j][i] = pdf_value * self.alpha[i]

            denominator = np.sum(probability, axis=1)

            # Computing the posterior probability
            for i in range(X.shape[0]):
                probability[i] = probability[i] / denominator[i]

            return probability

        # Computing the Log Likelihood
```

```

def compute_total_likelihood(self,X):
    probability = np.zeros((X.shape[0],self.K))
    for i in range(self.K):
        normal_fn = multivariate_normal(mean=self.means[i],cov=self.covs[i])
        for j in range(X.shape[0]):
            pdf_value = normal_fn.pdf(X[j])
            probability[j][i] = pdf_value* self.alpha[i]

    denominator = np.sum(probability,axis=1)
    for i in range(X.shape[0]):
        probability[i] = probability[i]/denominator[i]

    for i in range(X.shape[0]):
        for j in range(self.K):
            probability[i][j] = probability[i][j]*self.alpha[j]

    summation_on_mixture = np.log(np.sum(probability,axis=1))
    likelihood= np.sum(summation_on_mixture,axis=0)
    return likelihood

# Expectation step in EM Algorithm
def E_step(self):
    self.likelihoods = self.compute_posterior(self.data)

# Update step in EM Algorithm
def M_step(self):

    # Compute the new values of weights assigned for each gaussian
    new_alpha = {}
    mean_alpha = np.mean(self.likelihoods,axis=0)
    for i in range(self.K):
        new_alpha[i] = mean_alpha[i]

    denominator = np.sum(self.likelihoods,axis=0)

    # Compute the new mean vectors for the next iteration
    new_centers= {}
    # computing the terms in the numerator
    for i in range(self.K):
        logits = self.likelihoods[:,i]
        sum_vector = None
        for j in range(self.data.shape[0]):
            if sum_vector is None:
                sum_vector = logits[j]*self.data[j]
            else:
                sum_vector += logits[j]* self.data[j]

```

```

        new_centers[i]=sum_vector
    for i in range(self.K):
        new_centers[i] = new_centers[i]/denominator[i]

    # Compute the new covariance matrix for the next iteration

    new_covs = {}
    # computing the terms in the numerator
    for i in range(self.K):
        logits = self.likelihoods[:,i]
        sum_matrix = None
        for j in range(self.data.shape[0]):
            if sum_matrix is None:
                vec = (self.data[j]-new_centers[i]).reshape((-1,1))
                sum_matrix = logits[j]*(np.matmul(vec,np.transpose(vec)))
            else:
                sum_matrix += logits[j]*(np.matmul(np.transpose(vec),vec))

        new_covs[i]=sum_matrix

    for i in range(self.K):
        new_covs[i] = new_covs[i]/denominator[i]
        # Stripping of Non Diagonal Entries
        if(self.clip_nondiag_cov):
            diagonal_elements =new_covs[i].diagonal()
            new_covariance_matrix = np.eye(new_covs[i].shape[0])
            for j in range(self.K):
                new_covariance_matrix[j][j] = diagonal_elements[j]
            new_covs[i]= new_covariance_matrix

        if(np.linalg.matrix_rank(new_covs[i])!= 32):
            while(np.linalg.matrix_rank(new_covs[i])!= 32):
                new_covs[i] = new_covs[i] + np.eye(32)*1

    # Update for next iteration
    self.alpha = new_alpha
    self.centers= new_centers
    self.covs = new_covs

# The training Loop in the EM Algorithm
def train(self):
    print("GMM Training")
    for iter in range(1,self.num_iterations+1):

        loglikelihoods_value = self.compute_total_likelihood(self.data)

```

```

        print("Iteration" + str(iter) + " " + str(loglikelihoods_value))
        self.plot_points.append(loglikelihoods_value)
        self.E_step()
        self.M_step()

# Plotting the log likelihood as a function of EM Algorithm Iteration
def likelihood_fn(self):
    plt.xlabel("Iterations count")
    plt.ylabel("Log Likelihood")
    for i in range(len(self.plot_points)):
        plt.scatter(i, self.plot_points[i], c='r')
    plt.show()

```

## 1.1 Preparing the Datasets for Training and Testing

```

[5]: # Prepare the datasets
speech_train = Spectrogram("speech_music_classification/train/speech")
music_train = Spectrogram("speech_music_classification/train/music")

speech_train.compute_data_matrix()
music_train.compute_data_matrix()

```

```

[6]: speech_test = Spectrogram("speech_music_classification/train/speech")
music_test = Spectrogram("speech_music_classification/train/music")

```

```

[7]: # Helper function for checking accuracy
def accuracy_test(gmm_speech, gmm_music, dataset_type='speech'):

    directory = 'speech_music_classification/test'
    accuracy = 0
    total = 0
    spectrogram_class = Spectrogram()
    for file in os.listdir(directory):

        true_label = 0 if dataset_type in file else 1
        total+=1
        sample_rate, wav_file = wavfile.read(directory+"/"+file)
        spectrogram = spectrogram_class.compute_spectrogram(wav_file).
        ↪ reshape((32,-1))

        # Compute probability
        prob1 = gmm_speech.compute_posterior(spectrogram.transpose())
        prob2 = gmm_music.compute_posterior(spectrogram.transpose())

        # Holds the probability of each frame being belonging to a class
        prob1_average = np.zeros((prob1.shape[0],1))

```



```

prob2_average = np.zeros((prob2.shape[0],1))

for i in range(prob2.shape[0]):
    prob_speech = 0
    prob_music = 0
    for j in range(gmm_speech.K):
        prob_speech = prob_speech + prob1[i][j]*gmm_speech.alpha[j]
        prob_music = prob_music + prob2[i][j]*gmm_music.alpha[j]
    prob1_average[i] = prob_speech
    prob2_average[i] = prob_music

# Average over the probabilities of all the frames in audio file
prob_speech = np.mean(prob1_average,axis=0)
prob_music = np.mean(prob2_average,axis=0)

# Assigning to the argmax the predicted label
predicted_class = 0 if prob_speech >= prob_music else 1
if(predicted_class == true_label):
    accuracy+=1

print("Correctly classified = " + str(accuracy))
print("Total files = " + str(total))
accuracy = float(accuracy)/total
print("Accuracy of the GMM Predictions on test data = " + str(accuracy*100))

```

```

[8]: # Setting the parameters for Iteration count
kmeans_iteration_limit = 6
gmm_iteration_count = 15

```

## 2 2 Mixture Gaussian Diagonal Covariance

```

[9]: k_i_speech = KMeans(2,speech_train.data_matrix.
    ↪ transpose(),kmeans_iteration_limit,remove_non_diagonal=True)
k_i_speech.compute_means()
k_i_speech.compute_covariance()

```

```

[10]: gmm_i_speech = GMM(2,speech_train.data_matrix.
    ↪ transpose(),True,gmm_iteration_count)
gmm_i_speech.initialize(k_i_speech.centers,k_i_speech.covs)
gmm_i_speech.train()
gmm_i_speech.likelihood_fn()

```

GMM Training

```


Iteration1 -83122.20989274865
Iteration2 -83106.11487290719
Iteration3 -83095.72079375514

```

```

Iteration4 -83094.90861610003
Iteration5 -83094.83429937912
Iteration6 -83094.82663973875
Iteration7 -83094.82583955045
Iteration8 -83094.82575583678
Iteration9 -83094.8257470775
Iteration10 -83094.82574616118
Iteration11 -83094.82574606537
Iteration12 -83094.82574605512
Iteration13 -83094.82574605382
Iteration14 -83094.82574605383
Iteration15 -83094.8257460538

```



```

[11]: k_i_music = KMeans(2,music_train.data_matrix.
      ↳ transpose(),kmeans_iteration_limit,remove_non_diagonal=True)
      k_i_music.compute_means()
      k_i_music.compute_covariance()

```

```

[12]: gmm_i_music = GMM(2,music_train.data_matrix.transpose(),True,gmm_iteration_count)
      gmm_i_music.initialize(k_i_music.centers,k_i_music.covs)
      gmm_i_music.train()
      gmm_i_music.likelihood_fn()

```


GMM Training

```
Iteration1 -83122.20989274865
```

```

Iteration2 -82803.26901465219
Iteration3 -82805.32534540647
Iteration4 -82805.37531846325
Iteration5 -82805.38000960398
Iteration6 -82805.38045386915
Iteration7 -82805.38049593879
Iteration8 -82805.38049992235
Iteration9 -82805.38050029961
Iteration10 -82805.38050033535
Iteration11 -82805.38050033864
Iteration12 -82805.38050033872
Iteration13 -82805.38050033874
Iteration14 -82805.38050033877
Iteration15 -82805.3805003388

```



```
[13]: accuracy_test(gmm_i_speech,gmm_i_music)
```

Correctly classified = 40

Total files = 48

Accuracy of the GMM Predictions on test data = 83.33333333333334


### 3 2 mixture Gaussian Full Covariance

```
[14]: k_ii_speech = KMeans(2,speech_train.data_matrix.  
    ↪transpose()),kmeans_iteration_limit,remove_non_diagonal=False)  
k_ii_speech.compute_means()  
k_ii_speech.compute_covariance()
```

```
[15]: gmm_ii_speech = GMM(2,speech_train.data_matrix.  
    ↪transpose()),False,gmm_iteration_count)  
gmm_ii_speech.initialize(k_ii_speech.centers,k_ii_speech.covs)  
gmm_ii_speech.train()  
gmm_ii_speech.likelihood_fn()
```

*GMM Training*

```
Iteration1 -83122.20989274865  
Iteration2 -83120.99862376199  
Iteration3 -83120.36364427106  
Iteration4 -83120.26652164238  
Iteration5 -83120.25933612559  
Iteration6 -83120.25869087211  
Iteration7 -83120.25863161906  
Iteration8 -83120.25862616638  
Iteration9 -83120.25862566443  
Iteration10 -83120.25862561827  
Iteration11 -83120.25862561405  
Iteration12 -83120.25862561371  
Iteration13 -83120.25862561355  
Iteration14 -83120.25862561351  
Iteration15 -83120.25862561355
```




```
[16]: k_ii_music = KMeans(2,music_train.data_matrix.  
      ↪transpose(),kmeans_iteration_limit,remove_non_diagonal=False)  
      k_ii_music.compute_means()  
      k_ii_music.compute_covariance()
```

```
[17]: gmm_ii_music = GMM(2,music_train.data_matrix.  
      ↪transpose(),False,gmm_iteration_count)  
      gmm_ii_music.initialize(k_ii_music.centers,k_ii_music.covs)  
      gmm_ii_music.train()  
      gmm_ii_music.likelihood_fn()
```

#### GMM Training

```
Iteration1 -83122.20989274865  
Iteration2 -83048.51079647487  
Iteration3 -83056.33387379738  
Iteration4 -83057.56959489791  
Iteration5 -83057.7237151911  
Iteration6 -83057.73445543497  
Iteration7 -83057.73518287214  
Iteration8 -83057.73523204833  
Iteration9 -83057.73523537233  
Iteration10 -83057.73523559688  
Iteration11 -83057.73523561237  
Iteration12 -83057.73523561332
```

Iteration13 -83057.73523561323  
 Iteration14 -83057.73523561348  
 Iteration15 -83057.73523561348



```
[18]: accuracy_test(gmm_ii_speech,gmm_ii_music)
```

Correctly classified = 27

Total files = 48

Accuracy of the GMM Predictions on test data = 56.25

## 4 5 mixture Gaussian Diagonal Covariance

```
[19]: k_iii_speech = KMeans(5,speech_train.data_matrix.
    ↳ transpose(),kmeans_iteration_limit,remove_non_diagonal=True)
k_iii_speech.compute_means()
k_iii_speech.compute_covariance()
```

```
[20]: gmm_iii_speech = GMM(5,speech_train.data_matrix.
    ↳ transpose(),True,gmm_iteration_count)
gmm_iii_speech.initialize(k_iii_speech.centers,k_iii_speech.covs)
gmm_iii_speech.train()
gmm_iii_speech.likelihood_fn()
```


GMM Training

Iteration1 -193003.79445909726

```

Iteration2 -192525.1244459133
Iteration3 -192336.18078816903
Iteration4 -192346.62654893086
Iteration5 -192349.47876707264
Iteration6 -192350.15011428678
Iteration7 -192350.30620427165
Iteration8 -192350.34274367883
Iteration9 -192350.3513586448
Iteration10 -192350.35339957464
Iteration11 -192350.353884443
Iteration12 -192350.35399981032
Iteration13 -192350.3540272816
Iteration14 -192350.35403382572
Iteration15 -192350.35403538527

```



```

[21]: k_iii_music = KMeans(5,music_train.data_matrix.
      ↪ transpose(),kmeans_iteration_limit,remove_non_diagonal=True)
      k_iii_music.compute_means()
      k_iii_music.compute_covariance()

```


```

[22]: gmm_iii_music = GMM(5,music_train.data_matrix.
      ↪ transpose(),True,gmm_iteration_count)
      gmm_iii_music.initialize(k_iii_music.centers,k_iii_music.covs)
      gmm_iii_music.train()
      gmm_iii_music.likelihood_fn()

```

### *GMM Training*

```
Iteration1 -193003.79445909726
Iteration2 -191570.78945734003
Iteration3 -191404.44853250636
Iteration4 -191403.55549843112
Iteration5 -191403.52462655108
Iteration6 -191403.5253766211
Iteration7 -191403.5258871991
Iteration8 -191403.52598721322
Iteration9 -191403.52600271394
Iteration10 -191403.52600489388
Iteration11 -191403.52600518553
Iteration12 -191403.52600522415
Iteration13 -191403.5260052294
Iteration14 -191403.52600522953
Iteration15 -191403.52600522962
```



```
[23]: accuracy_test(gmm_iii_speech,gmm_iii_music)
```

*Correctly classified = 32*

*Total files = 48*

*Accuracy of the GMM Predictions on test data = 66.66666666666666*




## 5 5 mixture Gaussian Full Covariance

```
[24]: k_iv_speech = KMeans(5,speech_train.data_matrix.  
    ↪transpose()),kmeans_iteration_limit,remove_non_diagonal=False)  
k_iv_speech.compute_means()  
k_iv_speech.compute_covariance()
```

```
[25]: gmm_iv_speech = GMM(5,speech_train.data_matrix.  
    ↪transpose()),False,gmm_iteration_count)  
gmm_iv_speech.initialize(k_iv_speech.centers,k_iv_speech.covs)  
gmm_iv_speech.train()  
gmm_iv_speech.likelihood_fn()
```

*GMM Training*

```
Iteration1 -193003.79445909726  
Iteration2 -189338.5541192052  
Iteration3 -187285.63559388215  
Iteration4 -186584.5634945222  
Iteration5 -186013.42739394872  
Iteration6 -185572.08442673727  
Iteration7 -185246.28373146942  
Iteration8 -185011.74958000946  
Iteration9 -184845.27095740603  
Iteration10 -184728.07022512288  
Iteration11 -184645.977114116  
Iteration12 -184588.65922643358  
Iteration13 -184548.7229890435  
Iteration14 -184520.93589449147  
Iteration15 -184501.61997664248
```




```
[26]: k_iv_music = KMeans(5,music_train.data_matrix.
      ↪ transpose(),kmeans_iteration_limit,remove_non_diagonal=False)
      k_iv_music.compute_means()
      k_iv_music.compute_covariance()
```

```
[27]: gmm_iv_music = GMM(5,music_train.data_matrix.
      ↪ transpose(),False,gmm_iteration_count)
      gmm_iv_music.initialize(k_iv_music.centers,k_iv_music.covs)
      gmm_iv_music.train()
      gmm_iv_music.likelihood_fn()
```

#### GMM Training

```
Iteration1 -193003.79445909726
Iteration2 -184303.9517579474
Iteration3 -175568.0794464079
Iteration4 -173155.68166448345
Iteration5 -172127.8585550698
Iteration6 -171643.36049752237
Iteration7 -171402.86401053195
Iteration8 -171280.27003159418
Iteration9 -171216.9094585535
Iteration10 -171183.92621749875
Iteration11 -171166.69159167344
Iteration12 -171157.66826796025
Iteration13 -171152.9391509156
```

Iteration14 -171150.4592804641  
 Iteration15 -171149.15850776382



[28]: `accuracy_test(gmm_iv_speech,gmm_iv_music)`

*Correctly classified = 24*

*Total files = 48*

*Accuracy of the GMM Predictions on test data = 50.0*

## 6 Summary

	<i>Diagonal Covariance</i>	<i>Full Covariance</i>
<i>2 Mixture</i>	<i>83.34</i>	<i>56.25</i>
<i>5 Mixture</i>	<i>66.66</i>	<i>50.0</i>

### 6.1 Conclusions

- *The error rate increased on increasing the number of gaussians*
- *Fixing the number of mixture components, the Diagonal Covariance mixture models perform better than the full covariance models*

[ ]: